

## Architectural Soundness

The refactoring plan's approach is **architecturally sound** and follows solid design principles. By making `display_mode` a **derived property** (computed from `_selected_curves` and `_show_all_curves`), the design avoids duplicating state and guarantees consistency at all times <sup>1</sup>. This aligns with best practices like **single source of truth** for application state – all selection logic is centralized in `ApplicationState` instead of scattered across UI panels and widgets <sup>2</sup> <sup>3</sup>. Notably, the plan explicitly embraces these principles: *"derived state must be computed, never stored"* and *"UI components are input devices, not state owners"* <sup>4</sup>. In practice, this means UI elements (like the tracking panel or curve widget) simply feed user input into `ApplicationState` or read from it, rather than each maintaining their own mutable selection state. This separation of concerns is a recognized best practice for reliability and maintainability. It eliminates the previous situation where three different components held selection info that had to be manually kept in sync <sup>2</sup>. Overall, deriving `display_mode` from two authoritative fields ensures **there is no chance of mismatch** between what curves are selected vs. what mode is displayed, since the mode is always calculated on demand from the current inputs <sup>5</sup>. Removing mutable state from the UI (making the widget's `display_mode` read-only and driven by the central model) is indeed a best practice – it creates a clear **unidirectional data flow** (state → UI) and prevents UI and model from drifting out of sync. The plan's architecture is consistent with modern state management patterns (similar to MVC/MVVM or Redux-style stores) that improve consistency and testability. In short, centralizing selection state in `ApplicationState` with computed outputs is a sound approach that resolves the prior fragmented design and adheres to the "single source of truth" principle <sup>3</sup> <sup>5</sup>.

## Severe Risks / Edge Cases

The plan appears to have addressed the major bugs and race conditions from the old design (e.g. the bug where multi-selecting curves only showed one, and circular update loops on selection) <sup>6</sup> <sup>5</sup>. After refactoring, **no obvious race conditions remain** since the selection updates flow in one direction (UI input → `ApplicationState` → UI refresh) with no cyclical hand-backs. The removal of the panel→controller→panel feedback loop (which previously caused re-entrancy) is explicitly noted in the plan <sup>7</sup>, so that particular race condition is resolved. That said, there are a few things to watch for:

- **Outdated Usage of `display_mode`:** A breaking change is that code can no longer set `widget.display_mode` directly. If any legacy code or plugin still attempts to do so, it will either be a no-op or raise an error (since the setter is removed) <sup>8</sup>. This could lead to runtime errors or unexpected UI behavior if such code was missed in the refactor. The plan flags this as a risk ("⚠ Breaking change for code that sets `display_mode` directly" <sup>8</sup>), so it's crucial that **all** references were updated to use `ApplicationState` methods. A thorough search-and-replace (which the plan indicates was done in Phase 5) mitigates this, but it's a point of vigilance during code review and testing.
- **Mixed State or Partial Updates:** In any refactor that centralizes state, there's a risk that some component might temporarily hold an outdated view of state if not properly hooked up. For example, if a UI element doesn't listen to the new `selection_state_changed` signal, it might not update when selection changes. The plan addresses this by connecting change signals to trigger UI updates (e.g. ensuring the curve widget repaints on selection changes) <sup>9</sup> <sup>10</sup>. The use of `begin_batch()/end_batch()` when loading multiple curves is also important to avoid

intermediate inconsistent states or “signal storms” during bulk updates <sup>11</sup> <sup>12</sup>. As long as those batch wrappers and signal connections are implemented as planned, the risk of partial update glitches is low. It's worth double-checking that **every place** that modifies selection or show-all state goes through `ApplicationState`, so no component is left updating its old flags in isolation.

- **Edge Case: Toggle Combinations:** The logic for computing `display_mode` covers the basic combinations (show-all overrides any selection, otherwise selection vs none) <sup>1</sup>. An edge scenario to confirm is when a user rapidly toggles the “Show all curves” checkbox while having multiple curves selected. By design, if `show_all_curves` is `True`, the mode is `ALL_VISIBLE` regardless of selection; when it turns `False` again, the mode will revert to `SELECTED` if a selection set is still present <sup>1</sup>. This is the intended behavior, and no state is lost in the toggle (the selected set remains stored in `ApplicationState` while show-all is on). Ensuring the UI correctly reflects this – e.g. the panel's checkbox and table selection remain consistent – is important. The provided integration tests suggest this scenario is handled (they set the checkbox and verify state directly <sup>13</sup>). As an extra caution, one might test toggling off “Show all” when nothing was selected (which should yield `ACTIVE_ONLY` mode) to confirm the system gracefully falls back to the active curve display. The design inherently covers that case (empty selection + unchecked = active only) <sup>14</sup> <sup>15</sup>, so it should behave correctly.
- **Developer Misuse:** A subtle risk going forward is if future maintainers do not fully understand the new pattern and accidentally introduce new state variables or bypass `ApplicationState`. The plan's documentation and the ADR emphasize that selection state is computed and should not be manually set <sup>16</sup> <sup>17</sup>. As long as this is followed, the architecture will hold up. Any attempt to reintroduce ad-hoc flags in a UI widget (for example, someone adding another “selected” flag in a new component without integrating with `ApplicationState`) would reintroduce inconsistency. This is less a flaw in the plan and more a general caution: the team should enforce via code reviews that **all selection-related logic uses the new single source of truth**, and perhaps deprecate or remove any old APIs that allowed direct widget state mutation.

Overall, no severe design flaws are evident after the refactor – the main edge cases (multiple selection, toggling modes, no-selection state) are explicitly handled by the computed logic. The biggest risk is ensuring **complete adoption** of the new system so that no old pathways linger. With the plan's thorough phase-by-phase migration, this risk is minimized, but regression testing should specifically watch for any place a curve fails to display or a signal isn't hooked up (since those would indicate a missed integration of the new state model).

## Maintainability and Simplification

This plan greatly **simplifies the architecture** and should improve long-term maintainability. Removing duplicated state across UI components means there's now “one place to check selection state” and one place to update it <sup>5</sup>. The benefits – no sync bugs, easier debugging, cleaner design – are already noted in the plan <sup>18</sup>. There are a few areas to consider for even more clarity and future simplification:

- **Deprecate Legacy Signals/Fields:** The tracking panel's `display_mode_changed` signal and the widget's `selected_curve_names` list are kept for now to preserve backward compatibility <sup>19</sup> <sup>20</sup>, but these are essentially redundant in the new architecture. The plan wisely suggests marking the old signal as deprecated and eventually removing it in favor of the new `selection_state_changed` signal <sup>21</sup>. Doing this cleanup in a later release will reduce confusion (developers won't wonder which signal to use) and trim away any remaining vestiges

of the old state-handling. Similarly, any UI properties that are now just mirrors of `ApplicationState` (like `CurveViewWidget.selected_curve_names` which is updated only for backward compatibility <sup>22</sup> <sup>23</sup> ) could be refactored out once external code no longer depends on them. The end goal should be to **eliminate any duplicate state representations** – right now they’re mostly read-only mirrors, which is fine, but in the long term the codebase will be simpler when everything references `ApplicationState` directly.

- **Simplify TrackingPointsPanel Logic:** The panel currently contains logic to determine the display mode based on the checkbox and selection ( `_determine_mode_from_checkbox()` ) <sup>24</sup> <sup>25</sup> . With the new system, the panel doesn’t really need to calculate the mode itself – it just needs to update the two inputs (selected curves and show-all flag) and let `ApplicationState` derive the mode. In the refactored version, the panel’s checkbox handler likely calls `app_state.set_show_all_curves(checked)` (and the selection UI already emits the list of selected curves). This means the panel’s own internal `_display_mode` tracking and related re-entrancy guard might be vestigial. The plan’s integration test confirms that toggling the checkbox updates the `ApplicationState` immediately <sup>26</sup> , which implies the panel is directly driving `show_all_curves` now. We recommend **removing any now-unused complexity** in `TrackingPointsPanel` – for example, if `_on_display_mode_checkbox_toggled` no longer needs to emit its own `display_mode_changed` (because the controller is updating state directly), that code can be simplified. Every reduction in branching or state in the panel will make it easier to maintain. In summary, as the new architecture settles, devs should audit the UI code for any conditional logic that was only needed to sync with the old system, and remove it to streamline the codebase.
- **Documentation and Clarity:** The plan does an excellent job of adding inline documentation explaining the new selection state architecture <sup>27</sup> <sup>28</sup> and even provides an Architecture Decision Record <sup>29</sup> <sup>30</sup> . This is hugely beneficial for maintainability – new contributors can understand the rationale and design easily. One suggestion is to ensure these explanations are also summarized in the main developer guide or code contribution guidelines (so that no one inadvertently reverts to the old pattern). The principle “*make illegal states unrepresentable*” is followed here by not allowing `display_mode` to be set independently <sup>31</sup> , and that should be highlighted in documentation to guide future design decisions. With comprehensive docs and a clear ADR, the team has set a good foundation for long-term clarity.
- **Future Enhancements:** The centralized approach opens the door to easier enhancements – for instance, if in the future there’s a need for a new display mode or a different way to filter curves, having all inputs in one state object makes it straightforward to implement. The team should find it easier now to extend functionality without touching multiple GUI classes. One area of possible simplification is the **rendering logic**: currently the code updates `widget.selected_curves_ordered` for rendering order <sup>32</sup> <sup>33</sup> . In the future, it may be cleaner to have the render layer query `ApplicationState` for the list of curves to display (ordered as needed) rather than maintaining an ordered list on the widget. This wasn’t fully addressed in the plan (likely to minimize immediate changes), but it’s something to consider refactoring later for consistency. That said, these are minor points – the plan as written already produces a much cleaner design with far fewer moving parts. The **maintainability win** from this refactor is clear: there’s less duplicated logic to get wrong, and one canonical place to fix bugs or add features related to selection.

In summary, the plan not only improves the current code but also includes steps for cleanup (Phase 7) and documentation, which are key for maintainability. The codebase will be easier to work with after

this refactor, and any remaining complexity (mostly for backwards compatibility during transition) can be retired in the near future to simplify even more <sup>34</sup> <sup>35</sup> .

## Test and Dependency Coverage

The refactoring plan is backed by a comprehensive set of verification steps and tests, which should give a high degree of confidence in its correctness. The authors have outlined unit tests for the new `ApplicationState` logic (ensuring `display_mode` computes correctly for various input combinations) <sup>36</sup> <sup>37</sup> and have integration tests that cover the end-to-end flow from UI interactions to rendering outcomes. Notably, they plan a new `test_selection_state_integration.py` suite that **simulates user actions** on the tracking panel and verifies that `ApplicationState` and the UI respond appropriately <sup>38</sup> <sup>39</sup> . For example, one integration test will add two curves, simulate selecting both in the panel, and then assert that `app_state.get_selected_curves()` contains those curves and `display_mode == DisplayMode.SELECTED` <sup>40</sup> - effectively reproducing the original bug scenario to ensure it's fixed. Another test toggles the "Show all" checkbox and checks that `app_state.show_all_curves` becomes True and `display_mode` switches to ALL\_VISIBLE <sup>26</sup> . This coverage directly targets the core refactor goals and race conditions, which is excellent.

In addition to new tests, the plan calls for updating existing test suites (multi-point selection tests, render state tests, etc.) to use the new APIs <sup>41</sup> <sup>42</sup> . All places where tests used to set `widget.display_mode` or rely on the old selection flags will be revised to go through `ApplicationState` . This ensures that **existing behavior is not regressed** - the tests will fail if anything in the selection or rendering logic changed unintentionally. The plan even specifies aiming for **100% coverage** of the new selection-state methods <sup>43</sup> , which indicates a thorough approach. Type checking (via `./bpr` ) is also included in verification to catch any type or interface mismatches early <sup>44</sup> <sup>45</sup> .

One area to confirm is whether tests cover all edge cases of interest. The provided scenarios are strong (multiple selection, toggling show-all on and off, clearing selection). It would be wise to also test a scenario like "toggle show-all off when no curves are explicitly selected" (which should result in `ACTIVE_ONLY` mode). This might be implicitly covered by tests that clear the selection set and check the mode resets to active-only <sup>46</sup> <sup>47</sup> , but an integration test via the UI (e.g., deselecting all in the panel if possible) could be considered. Additionally, since the plan retains some legacy pathways for now, it could be useful to have a test ensuring that the deprecated `display_mode_changed` signal (if it's still emitted in the panel) does not cause any adverse effects - essentially that nothing significant is listening to it anymore. However, given that the controller has been changed to use `ApplicationState` signals instead, this is probably moot.

The **dependency coverage** seems adequate: the plan accounts for updating the controller, panel, widget, manager, and store in tandem, and the tests span from low-level state computation to high-level UI integration. The manual test script provided in the plan's verification steps <sup>48</sup> <sup>49</sup> is a good final sanity check - it uses the actual GUI components (`MainWindow`, `RenderState.compute` ) to mimic the user selecting two curves and asserts that both curves are visible in the rendered state <sup>50</sup> . This kind of end-to-end test (even if run manually) confirms that all parts of the system - from the panel's signals to the rendering logic - are working together under the new scheme.

In conclusion, the test coverage is **very thorough**. The combination of unit tests (for internal state logic), updated existing tests (to catch regressions in related functionality), and new integration tests (for the full user-level behavior) provides confidence that the refactoring will maintain correctness. No significant test scenario appears overlooked in the plan. If anything, the team might consider

automating the formerly manual GUI test if possible (the new integration tests go a long way toward that by instantiating widgets in a headless mode via `QApplication` <sup>51</sup> <sup>52</sup> ). With all tests passing and 100% of the new code exercised, the risk of uncaught issues is low. The careful phased approach – verifying at each step – means any bug introduced would be detected early in the process. Overall, the verification strategy is adequate and well thought out, giving a green light that this refactor can be merged with confidence in its correctness.

---

1 2 3 4 5 6 7 8 9 10 11 12 13 16 17 18 19 20 21 22 23 26 27 28 29 30 31 32 33  
34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52

SELECTION\_STATE\_REFACTORING\_DETAILED\_PLAN.md

file:///file-JqBFgRoiaL5oNH1K6hmEWf

14 15 24 25 tracking\_points\_panel.py

[https://github.com/semmlerino/CurveEditor/blob/70463c48f1cff6d79b4ac85122989e47e62211e5/ui/](https://github.com/semmlerino/CurveEditor/blob/70463c48f1cff6d79b4ac85122989e47e62211e5/ui/tracking_points_panel.py)

tracking\_points\_panel.py